

Advanced Hybrid CNN-Transformer Architecture for Early Lung Nodule Detection in CT scans

Submitted by:

Siddharth Koppisetty (CS22B1022)

Kolli Tanish (EE22B1017)

SOURCE CODE AND OUTPUTS:

Block 1

```
import os

import numpy as np

import pandas as pd

import SimpleITK as sitk

import matplotlib.pyplot as plt

from skimage import exposure

import pywt

from scipy import ndimage

import tensorflow as tf

from tensorflow.keras.models import load_model

import random

import cv2

from sklearn.model_selection import train_test_split

# Define paths

BASE_DIR = '/workspace/deep_learning_luna16/luna16_dataset' # Update with your path

ANNOTATIONS_PATH = os.path.join(BASE_DIR, 'annotations.csv')

CANDIDATES_PATH = os.path.join(BASE_DIR, 'candidates.csv')

SEGMENTATION_PATH = os.path.join(BASE_DIR, 'seg-lungs-LUNA16/seg-lungs-LUNA16')

SUBSET_PATHS = [os.path.join(BASE_DIR, f'subset{i}/subset{i}') for i in range(5)]

# Check if paths exist
```

```

print(f"Annotations file exists: {os.path.exists(ANNOTATIONS_PATH)}")
print(f"Candidates file exists: {os.path.exists(CANDIDATES_PATH)}")
print(f"Segmentation directory exists: {os.path.exists(SEGMENTATION_PATH)}")
for i, path in enumerate(SUBSET_PATHS):
    print(f"Subset{i} directory exists: {os.path.exists(path)}")

```

```

def load_scan(path):
    """Load CT scan from .mhd file"""
    itk_img = sitk.ReadImage(path)
    img_array = sitk.GetArrayFromImage(itk_img)
    origin = np.array(itk_img.GetOrigin())
    spacing = np.array(itk_img.GetSpacing())
    return img_array, origin, spacing, itk_img

```

```

def normalize_hu(image, min_bound=-1000, max_bound=400):
    """Normalize HU values to [0,1] range with lung window"""
    image = np.clip(image, min_bound, max_bound)
    image = (image - min_bound) / (max_bound - min_bound)
    return image

```

```

def standardize_spacing(image, spacing, new_spacing=[1.0, 1.0, 1.0]):
    """Standardize the spacing of the CT scan"""
    resize_factor = spacing / new_spacing
    new_shape = np.round(image.shape * resize_factor)
    real_resize_factor = new_shape / image.shape

    new_img = ndimage.zoom(image, real_resize_factor, mode='nearest')
    return new_img

```

```

def wavelet_enhancement(image, wavelet='db4', level=2, threshold_factor=1.0):
    """Enhance the image using wavelet decomposition"""

```

```
enhanced = np.zeros_like(image)
```

```
for i in range(image.shape[0]):
```

```
    slice_2d = image[i]
```

```
    # Wavelet decomposition
```

```
    coeffs = pywt.wavedec2(slice_2d, wavelet, level=level)
```

```
    # Modify the detail coefficients (high frequency components)
```

```
    coeffs_modified = [coeffs[0]] # Keep approximation coefficients as-is
```

```
    for j in range(1, len(coeffs)):
```

```
        # Each coeffs[j] is a tuple (Horizontal, Vertical, Diagonal)
```

```
        cH, cV, cD = coeffs[j]
```

```
        threshold_H = threshold_factor * np.std(cH)
```

```
        threshold_V = threshold_factor * np.std(cV)
```

```
        threshold_D = threshold_factor * np.std(cD)
```

```
        cH = pywt.threshold(cH, threshold_H, mode='soft')
```

```
        cV = pywt.threshold(cV, threshold_V, mode='soft')
```

```
        cD = pywt.threshold(cD, threshold_D, mode='soft')
```

```
        coeffs_modified.append((cH, cV, cD)) # Repack into a tuple
```

```
    # Reconstruct the image
```

```
    enhanced_slice = pywt.waverec2(coeffs_modified, wavelet)
```

```
    # In case the reconstructed slice size is slightly off due to padding
```

```
    enhanced[i] = enhanced_slice[:slice_2d.shape[0], :slice_2d.shape[1]]
```

```
# Normalize to [0,1]
```

```
enhanced = (enhanced - enhanced.min()) / (enhanced.max() - enhanced.min() + 1e-8)
```

```
return enhanced
```

```
def get_nodule_info():
```

```
"""Get nodule information from annotations.csv"""
```

```
annotations = pd.read_csv(ANNOTATIONS_PATH)
```

```
nodules = {}
```

```
for _, row in annotations.iterrows():
```

```
    seriesuid = row['seriesuid']
```

```
    coord_x = row['coordX']
```

```
    coord_y = row['coordY']
```

```
    coord_z = row['coordZ']
```

```
    diameter_mm = row['diameter_mm']
```

```
    if seriesuid not in nodules:
```

```
        nodules[seriesuid] = []
```

```
    nodules[seriesuid].append({
```

```
        'coord_x': coord_x,
```

```
        'coord_y': coord_y,
```

```
        'coord_z': coord_z,
```

```
        'diameter_mm': diameter_mm
```

```
    })
```

```
return nodules
```

```
def world_to_voxel(world_coords, origin, spacing):
```

```
    """Convert world coordinates to voxel coordinates"""
```

```
    voxel_coords = np.round((world_coords - origin) / spacing).astype(int)
```

```
    return voxel_coords
```

```
def extract_patches(image, center, diameter, patch_size=64):
```

```
    """Extract patches centered around nodules"""
```

```
    radius = int(diameter / 2)
```

```
patch_radius = patch_size // 2
```

```
# Get bounds
```

```
z_min = max(0, center[0] - patch_radius)
```

```
y_min = max(0, center[1] - patch_radius)
```

```
x_min = max(0, center[2] - patch_radius)
```

```
z_max = min(image.shape[0], center[0] + patch_radius)
```

```
y_max = min(image.shape[1], center[1] + patch_radius)
```

```
x_max = min(image.shape[2], center[2] + patch_radius)
```

```
# Extract patch
```

```
patch = np.zeros((patch_size, patch_size, patch_size))
```

```
extract = image[z_min:z_max, y_min:y_max, x_min:x_max]
```

```
# Place extract in the patch (handling edge cases)
```

```
z_extract, y_extract, x_extract = extract.shape
```

```
patch[:z_extract, :y_extract, :x_extract] = extract
```

```
return patch
```

```
import numpy as np
```

```
import random
```

```
from scipy import ndimage
```

```
from skimage import exposure
```

```
def autoaugment_3d(patch):
```

```
    """Apply a simplified and safer version of AutoAugment to 3D patches."""
```

```
    augmented = patch.copy()
```

```
# Clip before starting (to handle any small negatives from preprocessing)
```

```
augmented = np.clip(augmented, 0, 1)
```

```
# Randomly select 1 to 3 augmentations
```

```
operations = random.sample(  
    ['rotate', 'flip', 'translate', 'noise', 'contrast', 'identity'],  
    k=random.randint(1, 3)  
)
```

```
for op in operations:
```

```
    if op == 'rotate':
```

```
        # Random rotation angle (small)
```

```
        angle = random.uniform(-20, 20)
```

```
        for i in range(augmented.shape[0]): # Rotate each 2D slice
```

```
            augmented[i] = ndimage.rotate(augmented[i], angle, reshape=False, mode='nearest')
```

```
    elif op == 'flip':
```

```
        # Random flip along any axis
```

```
        flip_axis = random.randint(0, 2)
```

```
        augmented = np.flip(augmented, axis=flip_axis)
```

```
    elif op == 'translate':
```

```
        # Random shift along each axis
```

```
        shift = [random.randint(-5, 5) for _ in range(3)]
```

```
        augmented = ndimage.shift(augmented, shift, mode='nearest')
```

```
    elif op == 'noise':
```

```
        # Add Gaussian noise
```

```
        noise = np.random.normal(0, 0.02, size=augmented.shape)
```

```
        augmented = np.clip(augmented + noise, 0, 1)
```

```
    elif op == 'contrast':
```

```

        # Random contrast adjustment

        # Important: Clip to [0,1] before contrast adjustment
        augmented = np.clip(augmented, 0, 1)
        gamma = random.uniform(0.8, 1.2)
        augmented = exposure.adjust_gamma(augmented, gamma)

    elif op == 'identity':
        # Do nothing (keep the original patch)
        pass

    # Final safety: make sure final augmented is in [0,1]
    augmented = np.clip(augmented, 0, 1)

    return augmented

def medmix_augmentation(patch, mask, alpha=0.5):
    """Apply MedMix augmentation to patches"""
    # For simplicity, we'll assume the mask is binary (1s where nodule is)
    # In a real implementation, you'd have accurate segmentation masks

    # Get another random patch (would be from another sample in practice)
    # Here we'll just use a slightly modified version of the original
    other_patch = patch * 0.9 + np.random.normal(0, 0.1, patch.shape)

    # MedMix: Keep the region of interest from the original
    # and mix the background with the other patch
    mixed_patch = mask * patch + (1 - mask) * (alpha * patch + (1 - alpha) * other_patch)

    return mixed_patch

def prepare_dataset():

```

```

"""Main function to prepare the dataset"""

# Get nodule information
nodules_info = get_nodule_info()

X_positive = [] # Patches with nodules
X_negative = [] # Patches without nodules

# Process each subset
for subset_dir in SUBSET_PATHS:
    subset_name = os.path.basename(subset_dir)
    print(f"Processing {subset_name}...")

    # Get all .mhd files in the subset
    for root, _, files in os.walk(subset_dir):
        for file in files:
            if file.endswith('.mhd'):
                file_path = os.path.join(root, file)
                series_id = file.split('.')[0]

                # Check if this series has nodules
                if series_id in nodules_info:
                    # Load the scan
                    img_array, origin, spacing, _ = load_scan(file_path)

                    # Step 1: Preprocessing
                    # Load segmentation mask (if available)
                    seg_path = os.path.join(SEGMENTATION_PATH, f"{series_id}.npy")
                    if os.path.exists(seg_path):
                        lung_mask = np.load(seg_path)
                    else:
                        # If segmentation not available, use a simple threshold

```



```

lung_mask = (img_array > -400) & (img_array < 0)

# Apply lung mask and normalize
img_array = img_array * lung_mask
img_array = normalize_hu(img_array)

# Standardize spacing
img_array = standardize_spacing(img_array, spacing)

# Step 2: Enhancement using wavelet decomposition
enhanced_img = wavelet_enhancement(img_array)

# Extract positive patches (around nodules)
for nodule in nodules_info[series_id]:
    world_coords = np.array([nodule['coord_z'], nodule['coord_y'], nodule['coord_x']])
    voxel_coords = world_to_voxel(world_coords, origin, spacing)

    # Make sure coordinates are within image bounds
    if (0 <= voxel_coords[0] < img_array.shape[0] and
        0 <= voxel_coords[1] < img_array.shape[1] and
        0 <= voxel_coords[2] < img_array.shape[2]):

        # Extract patch around nodule
        patch = extract_patches(enhanced_img, voxel_coords,
                                nodule['diameter_mm'], patch_size=64)

        # Create a simple spherical mask for the nodule (approximate)
        mask = np.zeros_like(patch)
        center = np.array(patch.shape) // 2
        for z in range(patch.shape[0]):
            for y in range(patch.shape[1]):

```

```

        for x in range(patch.shape[2]):
            dist = np.sqrt(((z-center[0])**2 + (y-center[1])**2 + (x-center[2])**2))
            if dist <= nodule['diameter_mm'] / 2:
                mask[z, y, x] = 1

# Step 3: Augmentation

# Store original
X_positive.append(patch)

# Generate augmented versions
for _ in range(3): # Create 3 augmented samples per nodule
    aug_patch = autoaugment_3d(patch)
    X_positive.append(aug_patch)

# Apply MedMix augmentation
medmix_patch = medmix_augmentation(patch, mask)
X_positive.append(medmix_patch)

# Extract some negative patches (regions without nodules)
# This is a simplified approach - in practice you'd want to be more selective
for _ in range(min(len(nodules_info[series_id]) * 5, 20)):
    # Random position (away from nodules)
    z = random.randint(32, img_array.shape[0] - 32)
    y = random.randint(32, img_array.shape[1] - 32)
    x = random.randint(32, img_array.shape[2] - 32)

    # Make sure it's far from any nodule
    is_far = True
    for nodule in nodules_info[series_id]:
        world_coords = np.array([nodule['coord_z'], nodule['coord_y'], nodule['coord_x']])
        voxel_coords = world_to_voxel(world_coords, origin, spacing)

```

```

dist = np.sqrt((z - voxel_coords[0])**2 +
               (y - voxel_coords[1])**2 +
               (x - voxel_coords[2])**2)

if dist < 50: # Arbitrary distance threshold
    is_far = False
    break

if is_far:
    center = np.array([z, y, x])
    patch = extract_patches(enhanced_img, center, 20, patch_size=64)
    X_negative.append(patch)

# Convert to numpy arrays
X_positive = np.array(X_positive)
X_negative = np.array(X_negative)

# Create labels
y_positive = np.ones(len(X_positive))
y_negative = np.zeros(len(X_negative))

# Combine datasets
X = np.concatenate([X_positive, X_negative], axis=0)
y = np.concatenate([y_positive, y_negative], axis=0)

# Shuffle
indices = np.arange(len(X))
np.random.shuffle(indices)
X = X[indices]
y = y[indices]

```

```

# Split into train/val/test sets

X_train, X_temp, y_train, y_temp = train_test_split(X, y, test_size=0.3, random_state=42)
X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp, test_size=0.5, random_state=42)


print(f"Dataset prepared with {len(X_positive)} positive and {len(X_negative)} negative samples")
print(f"Train: {len(X_train)}, Validation: {len(X_val)}, Test: {len(X_test)}")


return X_train, X_val, X_test, y_train, y_val, y_test


def visualize_sample(image, title="Sample"):
    """Visualize middle slice of a 3D volume"""
    middle_slice = image[image.shape[0]//2]
    plt.figure(figsize=(10, 10))
    plt.imshow(middle_slice, cmap='gray')
    plt.title(title)
    plt.colorbar()
    plt.show()


if __name__ == "__main__":
    # Update the base directory path

    BASE_DIR = '/workspace/deep_learning_luna16/luna16_dataset' # Change this to your actual path


    # Prepare dataset

    X_train, X_val, X_test, y_train, y_val, y_test = prepare_dataset()


    # Visualize some samples

    pos_idx = np.where(y_train == 1)[0][0]
    neg_idx = np.where(y_train == 0)[0][0]

```


```
visualize_sample(X_train[pos_idx], "Positive Sample (with nodule)")  
visualize_sample(X_train[neg_idx], "Negative Sample (no nodule)")
```

```
# Save preprocessed data  
  
np.save('X_train.npy', X_train)  
np.save('X_val.npy', X_val)  
np.save('X_test.npy', X_test)  
np.save('y_train.npy', y_train)  
np.save('y_val.npy', y_val)  
np.save('y_test.npy', y_test)
```

```
print("Preprocessing, enhancement, and augmentation complete!")
```

OUTPUT OF BLOCK 1:

```
Annotations file exists: True  
Candidates file exists: True  
Segmentation directory exists: True  
Subset0 directory exists: True  
Subset1 directory exists: True  
Subset2 directory exists: True  
Subset3 directory exists: True  
Subset4 directory exists: True  
Processing subset0...  
Processing subset1...  
Processing subset2...  
Processing subset3...  
Processing subset4...  
Dataset prepared with 780 positive and 2831 negative samples  
Train: 2527, Validation: 542, Test: 542
```



BLOCK 2:

```
import numpy as np

import tensorflow as tf

from tensorflow.keras import layers, models, optimizers

from tensorflow.keras.callbacks import ModelCheckpoint, EarlyStopping, ReduceLROnPlateau

from tensorflow.keras.losses import BinaryCrossentropy

from tensorflow.keras import backend as K

import matplotlib.pyplot as plt

from sklearn.metrics import roc_curve, auc, confusion_matrix, precision_recall_curve

import pandas as pd

import seaborn as sns

import os


# For reproducibility
np.random.seed(42)
tf.random.set_seed(42)


def focal_loss(gamma=2., alpha=.25):
    """Focal loss for addressing class imbalance"""

    def focal_loss_fixed(y_true, y_pred):
        pt_1 = tf.where(tf.equal(y_true, 1), y_pred, tf.ones_like(y_pred))
```

```
pt_0 = tf.where(tf.equal(y_true, 0), y_pred, tf.zeros_like(y_pred))
```

```
epsilon = K.epsilon()
```

```
pt_1 = K.clip(pt_1, epsilon, 1. - epsilon)
```

```
pt_0 = K.clip(pt_0, epsilon, 1. - epsilon)
```

```
return -K.mean(alpha * K.pow(1. - pt_1, gamma) * K.log(pt_1)) \
```

```
      -K.mean((1 - alpha) * K.pow(pt_0, gamma) * K.log(1. - pt_0))
```

```
return focal_loss_fixed
```

```
from tensorflow.keras import layers
```

```
from tensorflow.keras import layers
```

```
def efficient_3d_block(x, filters, kernel_size=3, stride=1, expand_ratio=1, se_ratio=0.25,  
activation='swish'):
```

```
    shortcut = x
```

```
    input_channels = x.shape[-1]
```

```
    expanded_channels = input_channels * expand_ratio
```

```
    # Expansion phase
```

```
    if expand_ratio != 1:
```

```
        x = layers.Conv3D(expanded_channels, kernel_size=1, padding='same', use_bias=False)(x)
```

```
        x = layers.BatchNormalization()(x)
```

```
        x = layers.Activation(activation)(x)
```

```
    # "Depthwise" convolution using groups
```

```
    x = layers.Conv3D(  
        expanded_channels,  
        kernel_size=kernel_size,  
        strides=stride,
```

```

padding='same',
use_bias=False,
groups=expanded_channels
)(x)
x = layers.BatchNormalization()(x)
x = layers.Activation(activation)(x)

# Squeeze and Excitation
if se_ratio > 0:
    se_channels = max(1, int(input_channels * se_ratio))
    se = layers.GlobalAveragePooling3D()(x)
    se = layers.Reshape((1, 1, 1, expanded_channels))(se)
    se = layers.Conv3D(se_channels, kernel_size=1, activation=activation, padding='same')(se)
    se = layers.Conv3D(expanded_channels, kernel_size=1, activation='sigmoid', padding='same')(se)
    x = layers.Multiply()([x, se])

```

```

# Projection phase
x = layers.Conv3D(filters, kernel_size=1, padding='same', use_bias=False)(x)
x = layers.BatchNormalization()(x)

```

```

# Skip connection
if stride == 1 and input_channels == filters:
    x = layers.Add()([shortcut, x])

```

```

return x

```

```

def transformer_encoder_block(x, embed_dim, num_heads, ff_dim, dropout=0.1):

```

```

    """Transformer encoder block for global context modeling"""

```

```

    # Multi-head attention

```

```

    attention_output = layers.MultiHeadAttention(
        num_heads=num_heads, key_dim=embed_dim // num_heads)(x, x)

```



```

attention_output = layers.Dropout(dropout)(attention_output)
x1 = layers.Add()([x, attention_output])
x1 = layers.LayerNormalization(epsilon=1e-6)(x1)

```

Feed-forward network

```

ffn_output = layers.Dense(ff_dim, activation='relu')(x1)
ffn_output = layers.Dense(embed_dim)(ffn_output)
ffn_output = layers.Dropout(dropout)(ffn_output)
x2 = layers.Add()([x1, ffn_output])
out = layers.LayerNormalization(epsilon=1e-6)(x2)

```

```

return out

```

```

def build_model(input_shape=(64, 64, 64, 1), base_filters=32):

```

```

    """Build the hybrid CNN-Transformer model for lung nodule detection"""

```

```

    inputs = layers.Input(shape=input_shape)

```

Initial Conv

```

x = layers.Conv3D(base_filters, kernel_size=3, strides=1, padding='same')(inputs)
x = layers.BatchNormalization()(x)
x = layers.Activation('relu')(x)

```

Progressive complexity - EfficientNet-style blocks

Block 1

```

x = efficient_3d_block(x, base_filters)
x = layers.MaxPooling3D(pool_size=2)(x)

```

Block 2

```

x = efficient_3d_block(x, base_filters * 2)
x = efficient_3d_block(x, base_filters * 2)
x = layers.MaxPooling3D(pool_size=2)(x)

```

Block 3

```
x = efficient_3d_block(x, base_filters * 4)
```

```
x = efficient_3d_block(x, base_filters * 4)
```

```
x = layers.MaxPooling3D(pool_size=2)(x)
```

Block 4

```
x = efficient_3d_block(x, base_filters * 8)
```

```
x = efficient_3d_block(x, base_filters * 8)
```

Transformer for global context

Reshape for transformer

```
batch_size = tf.shape(x)[0]
```

```
h_dim = tf.shape(x)[1]
```

```
w_dim = tf.shape(x)[2]
```

```
d_dim = tf.shape(x)[3]
```

```
c_dim = tf.shape(x)[4]
```

Flatten spatial dimensions for transformer

```
x_resaped = tf.reshape(x, [batch_size, h_dim * w_dim * d_dim, c_dim])
```

Apply transformer blocks

```
transformer_dim = base_filters * 8
```

```
x_transformed = transformer_encoder_block(
```

```
    x_resaped, embed_dim=transformer_dim,
```

```
    num_heads=4, ff_dim=transformer_dim * 4, dropout=0.1
```

```
)
```

Global pooling after transformer

```
x_pooled = layers.GlobalAveragePooling1D()(x_transformed)
```

Classification head

```
x = layers.Dense(256, activation='relu')(x_pooled)
```

```
x = layers.Dropout(0.3)(x)
```

```
x = layers.Dense(64, activation='relu')(x)
```

```
x = layers.Dropout(0.2)(x)
```

Output with uncertainty (logits)

```
output = layers.Dense(1, activation='sigmoid')(x)
```

```
model = models.Model(inputs, output)
```

```
return model
```

```
def train_model(X_train, y_train, X_val, y_val, batch_size=16, epochs=50, lr=0.001):
```

```
    """Train the model with progressive learning strategy"""
```

Reshape data to include channel dimension if needed

```
if len(X_train.shape) == 4:
```

```
    X_train = np.expand_dims(X_train, axis=-1)
```

```
    X_val = np.expand_dims(X_val, axis=-1)
```

Build model

```
model = build_model(input_shape=X_train.shape[1:])
```

Compile with focal loss

```
model.compile(
```

```
    optimizer=optimizers.Adam(learning_rate=lr),
```

```
    loss=focal_loss(gamma=2.0, alpha=0.25),
```

```
    metrics=['accuracy', tf.keras.metrics.Precision(), tf.keras.metrics.Recall(),
```

```
            tf.keras.metrics.AUC()])
```

```
)
```

Setup callbacks

```
checkpoint = ModelCheckpoint(
    'best_model.h5', monitor='val_auc', verbose=1,
    save_best_only=True, mode='max'
)
```

```
early_stopping = EarlyStopping(
    monitor='val_auc', patience=10, verbose=1, mode='max',
    restore_best_weights=True
)
```

```
reduce_lr = ReduceLROnPlateau(
    monitor='val_loss', factor=0.5, patience=5, verbose=1,
    min_lr=1e-6
)
```

```
callbacks = [checkpoint, early_stopping, reduce_lr]
```

Progressive learning - start with small subset and gradually increase

Stage 1: 25% of data

```
print("Stage 1: Training with 25% of data...")
```

```
idx = np.random.choice(len(X_train), size=int(len(X_train)*0.25), replace=False)
```

```
history_1 = model.fit(
    X_train[idx], y_train[idx],
    batch_size=batch_size,
    epochs=int(epochs/3),
    validation_data=(X_val, y_val),
    callbacks=callbacks,
    verbose=1
)
```

Stage 2: 50% of data

```
print("Stage 2: Training with 50% of data...")

idx = np.random.choice(len(X_train), size=int(len(X_train)*0.5), replace=False)

history_2 = model.fit(
    X_train[idx], y_train[idx],
    batch_size=batch_size,
    epochs=int(epochs/3),
    validation_data=(X_val, y_val),
    callbacks=callbacks,
    verbose=1
)
```

Stage 3: 100% of data

```
print("Stage 3: Training with 100% of data...")

history_3 = model.fit(
    X_train, y_train,
    batch_size=batch_size,
    epochs=int(epochs/3),
    validation_data=(X_val, y_val),
    callbacks=callbacks,
    verbose=1
)
```

Combine histories

```
combined_history = {}

for key in history_1.history.keys():
    combined_history[key] = (
        history_1.history[key] +
        history_2.history[key] +
        history_3.history[key]
    )
```

```
return model, combined_history
```

```
def evaluate_model(model, X_test, y_test):
```

```
    """Evaluate model performance with clinical metrics"""
```

```
    # Get predictions
```

```
    y_pred_prob = model.predict(X_test)
```

```
    y_pred = (y_pred_prob > 0.5).astype(int).flatten()
```

```
    # Calculate ROC curve and AUC
```

```
    fpr, tpr, thresholds = roc_curve(y_test, y_pred_prob)
```

```
    roc_auc = auc(fpr, tpr)
```

```
    # Find clinical thresholds (high sensitivity)
```

```
    # Aim for > 95% sensitivity for cancer detection
```

```
    target_sensitivity = 0.95
```

```
    idx = np.argmax(tpr >= target_sensitivity)
```

```
    clinical_threshold = thresholds[idx]
```

```
    clinical_sensitivity = tpr[idx]
```

```
    clinical_specificity = 1 - fpr[idx]
```

```
    # Apply clinical threshold
```

```
    y_pred_clinical = (y_pred_prob > clinical_threshold).astype(int).flatten()
```

```
    # Calculate confusion matrix
```

```
    cm = confusion_matrix(y_test, y_pred_clinical)
```

```
    # False negative analysis
```

```
    false_negatives = np.where((y_test == 1) & (y_pred_clinical == 0))[0]
```

```
    # Print metrics
```

```
    print(f"AUC-ROC: {roc_auc:.4f}")
```

```

print(f"Clinical threshold: {clinical_threshold:.4f}")
print(f"Sensitivity at clinical threshold: {clinical_sensitivity:.4f}")
print(f"Specificity at clinical threshold: {clinical_specificity:.4f}")
print(f"Number of false negatives: {len(false_negatives)}")

# Plot ROC curve
plt.figure(figsize=(10, 8))
plt.plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve (AUC = {roc_auc:.4f})')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.scatter(fpr[idx], tpr[idx], marker='o', color='red',
            label=f"Clinical threshold (Sens={clinical_sensitivity:.2f}, Spec={clinical_specificity:.2f})")
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic')
plt.legend(loc="lower right")
plt.savefig('roc_curve.png')
plt.show()

# Plot confusion matrix
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=['No Nodule', 'Nodule'],
            yticklabels=['No Nodule', 'Nodule'])
plt.ylabel('True Label')
plt.xlabel('Predicted Label')
plt.title('Confusion Matrix at Clinical Threshold')
plt.savefig('confusion_matrix.png')
plt.show()

```

```
# Plot precision-recall curve
```

```
precision, recall, _ = precision_recall_curve(y_test, y_pred_prob)
```

```
plt.figure(figsize=(10, 8))
```

```
plt.plot(recall, precision, color='blue', lw=2)
```

```
plt.xlabel('Recall')
```

```
plt.ylabel('Precision')
```

```
plt.title('Precision-Recall Curve')
```

```
plt.savefig('precision_recall_curve.png')
```

```
plt.show()
```

```
return {
```

```
    'auc_roc': roc_auc,
```

```
    'clinical_threshold': clinical_threshold,
```

```
    'sensitivity': clinical_sensitivity,
```

```
    'specificity': clinical_specificity,
```

```
    'false_negatives': len(false_negatives)
```

```
}
```

```
def cross_dataset_validation(model, external_datasets):
```

```
    """Validate model on external datasets to test generalizability"""
```

```
    results = {}
```

```
    for dataset_name, (X, y) in external_datasets.items():
```

```
        print(f"\nEvaluating on external dataset: {dataset_name}")
```

```
        # Reshape if needed
```

```
        if len(X.shape) == 4:
```

```
            X = np.expand_dims(X, axis=-1)
```

```
        # Predict and evaluate
```

```
        metrics = evaluate_model(model, X, y)
```



```

    results[dataset_name] = metrics

# Compare performance across datasets
datasets = list(results.keys())
aucs = [results[d]['auc_roc'] for d in datasets]
sensitivities = [results[d]['sensitivity'] for d in datasets]
specificities = [results[d]['specificity'] for d in datasets]

# Plot comparison
plt.figure(figsize=(12, 6))
x = np.arange(len(datasets))
width = 0.25

plt.bar(x - width, aucs, width, label='AUC-ROC')
plt.bar(x, sensitivities, width, label='Sensitivity')
plt.bar(x + width, specificities, width, label='Specificity')

plt.ylabel('Score')
plt.title('Performance Across Datasets')
plt.xticks(x, datasets)
plt.legend()
plt.savefig('cross_dataset_performance.png')
plt.show()

return results

def plot_training_history(history):
    """Plot training metrics history"""
    plt.figure(figsize=(12, 5))

    # Plot loss

```

```
plt.subplot(1, 2, 1)
plt.plot(history['loss'], label='Training Loss')
plt.plot(history['val_loss'], label='Validation Loss')
plt.title('Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
```

Plot AUC

```
plt.subplot(1, 2, 2)
plt.plot(history['auc'], label='Training AUC')
plt.plot(history['val_auc'], label='Validation AUC')
plt.title('AUC')
plt.xlabel('Epoch')
plt.ylabel('AUC')
plt.legend()
```

```
plt.tight_layout()
plt.savefig('training_history.png')
plt.show()
```

```
def analyze_false_negatives(model, X_test, y_test, threshold=0.5):
```

```
    """Analyze false negative cases which are critical in cancer detection"""
```

```
    y_pred_prob = model.predict(X_test)
```

```
    y_pred = (y_pred_prob > threshold).astype(int).flatten()
```

Find false negatives

```
fn_indices = np.where((y_test == 1) & (y_pred == 0))[0]
```

```
if len(fn_indices) == 0:
```

```
    print("No false negatives found at this threshold.")
```

```

return

print(f"Analyzing {len(fn_indices)} false negative cases...")

# Get probabilities for false negatives
fn_probs = y_pred_prob[fn_indices].flatten()

# Plot probability distribution
plt.figure(figsize=(10, 6))
plt.hist(fn_probs, bins=20, alpha=0.7)
plt.axvline(threshold, color='red', linestyle='--', label=f'Threshold: {threshold}')
plt.title('Probability Distribution of False Negatives')
plt.xlabel('Prediction Probability')
plt.ylabel('Count')
plt.legend()
plt.savefig('false_negative_distribution.png')
plt.show()

# Visualize some false negative examples
num_examples = min(5, len(fn_indices))
plt.figure(figsize=(15, 3 * num_examples))

for i in range(num_examples):
    idx = fn_indices[i]
    prob = y_pred_prob[idx][0]

    # Get middle slice of the volume
    if len(X_test.shape) == 5: # Check if channel dimension exists
        middle_slice = X_test[idx, X_test.shape[1]//2, :, :, 0]
    else:
        middle_slice = X_test[idx, X_test.shape[1]//2, :, :]

```

```

plt.subplot(num_examples, 1, i+1)

plt.imshow(middle_slice, cmap='gray')

plt.title(f'False Negative #{i+1}: Prediction Probability = {prob:.4f}')

plt.colorbar()


plt.tight_layout()

plt.savefig('false_negative_examples.png')

plt.show()


return fn_indices, fn_probs


def save_results(metrics, history, output_dir='results'):
    """Save evaluation results to file"""

    if not os.path.exists(output_dir):
        os.makedirs(output_dir)


    # Save metrics as CSV
    metrics_df = pd.DataFrame([metrics])
    metrics_df.to_csv(os.path.join(output_dir, 'evaluation_metrics.csv'), index=False)


    # Save history as CSV
    history_df = pd.DataFrame(history)
    history_df.to_csv(os.path.join(output_dir, 'training_history.csv'), index=False)


    print(f"Results saved to {output_dir}")


if __name__ == "__main__":
    # Load preprocessed data
    X_train = np.load('X_train.npy')
    X_val = np.load('X_val.npy')

```

```

X_test = np.load('X_test.npy')
y_train = np.load('y_train.npy')
y_val = np.load('y_val.npy')
y_test = np.load('y_test.npy')

# Add channel dimension if needed
if len(X_train.shape) == 4: # (samples, depth, height, width)
    X_train = np.expand_dims(X_train, axis=-1)
    X_val = np.expand_dims(X_val, axis=-1)
    X_test = np.expand_dims(X_test, axis=-1)

print(f"Training data shape: {X_train.shape}")
print(f"Validation data shape: {X_val.shape}")
print(f"Test data shape: {X_test.shape}")

# Train the model
model, history = train_model(
    X_train, y_train,
    X_val, y_val,
    batch_size=8, # Smaller batch size for 3D data
    epochs=60,
    lr=0.0003
)

# Plot training history
plot_training_history(history)

# Evaluate on test set
print("\nEvaluating on test set...")
metrics = evaluate_model(model, X_test, y_test)

```

```
# Analyze false negatives (critical in cancer detection)

analyze_false_negatives(model, X_test, y_test, threshold=metrics['clinical_threshold'])
```

```
# Save results

save_results(metrics, history)
```

OUTPUT OF BLOCK 2:

BLOCK 3

```
import os

import numpy as np

import tensorflow as tf

import matplotlib.pyplot as plt

import matplotlib.animation as animation

from mpl_toolkits.axes_grid1 import make_axes_locatable

from tensorflow.keras import backend as K


# Define focal loss

def focal_loss(gamma=2., alpha=.25):

    """Focal loss for addressing class imbalance"""

    def focal_loss_fixed(y_true, y_pred):

        pt_1 = tf.where(tf.equal(y_true, 1), y_pred, tf.ones_like(y_pred))

        pt_0 = tf.where(tf.equal(y_true, 0), y_pred, tf.zeros_like(y_pred))

        epsilon = K.epsilon()

        pt_1 = K.clip(pt_1, epsilon, 1. - epsilon)

        pt_0 = K.clip(pt_0, epsilon, 1. - epsilon)

        return -K.mean(alpha * K.pow(1. - pt_1, gamma) * K.log(pt_1)) \

            -K.mean((1 - alpha) * K.pow(pt_0, gamma) * K.log(1. - pt_0))

    return focal_loss_fixed


# Visualization functions

def visualize_3d_volume(volume, title="3D Volume Visualization", save_path=None):

    if len(volume.shape) == 4:
```

```

    vol = volume[..., 0]
else:
    vol = volume
depth = vol.shape[0]
fig, ax = plt.subplots(figsize=(10, 8))
plt.tight_layout()
slice_idx = depth // 2
img = ax.imshow(vol[slice_idx], cmap='gray')
divider = make_axes_locatable(ax)
cax = divider.append_axes("right", size="5%", pad=0.05)
plt.colorbar(img, cax=cax)
title_obj = ax.set_title(f"{title} - Slice {slice_idx+1}/{depth}")
def update(frame):
    img.set_data(vol[frame])
    title_obj.set_text(f"{title} - Slice {frame+1}/{depth}")
    return [img, title_obj]
ani = animation.FuncAnimation(fig, update, frames=depth, interval=100, blit=True)
if save_path:
    os.makedirs(os.path.dirname(save_path), exist_ok=True)
    ani.save(save_path, writer='pillow', fps=10)
plt.close()
return ani

```

```

def visualize_prediction_comparison(volume, true_label, pred_prob, save_dir=None):
    fig, axes = plt.subplots(1, 3, figsize=(15, 5))
    if len(volume.shape) == 4:
        vol = volume[..., 0]
    else:
        vol = volume
    center_z = vol.shape[0] // 2
    center_y = vol.shape[1] // 2

```

```

center_x = vol.shape[2] // 2
axes[0].imshow(vol[center_z], cmap='gray')
axes[0].set_title(f"Axial (Z={center_z})")
axes[0].axis('off')
axes[1].imshow(vol[:, center_y], cmap='gray')
axes[1].set_title(f"Coronal (Y={center_y})")
axes[1].axis('off')
axes[2].imshow(vol[:, :, center_x], cmap='gray')
axes[2].set_title(f"Sagittal (X={center_x})")
axes[2].axis('off')

pred_text = f"True: {'Nodule' if true_label == 1 else 'No Nodule'}\n" \
            f"Pred: {pred_prob:.4f}\n" \
            f"Class: {'Nodule' if pred_prob > 0.5 else 'No Nodule'}"
fig.suptitle(pred_text, fontsize=14)
plt.tight_layout()
if save_dir:
    os.makedirs(save_dir, exist_ok=True)
    plt.savefig(os.path.join(save_dir, f"pred_comparison_{true_label}_{pred_prob:.2f}.png"))
plt.close()
return fig

```

```

def visualize_model_predictions(model, X_test, y_test, num_samples=5, save_dir=None):

```

```

    if not save_dir:
        save_dir = '3d_res'
    os.makedirs(save_dir, exist_ok=True)
    y_pred_prob = model.predict(X_test)
    # Get indices for TP, TN, FP, FN
    tp_indices = np.where((y_test == 1) & (y_pred_prob.flatten() > 0.5))[0]
    tn_indices = np.where((y_test == 0) & (y_pred_prob.flatten() <= 0.5))[0]
    fp_indices = np.where((y_test == 0) & (y_pred_prob.flatten() > 0.5))[0]
    fn_indices = np.where((y_test == 1) & (y_pred_prob.flatten() <= 0.5))[0]

```



```

fig, axes = plt.subplots(4, num_samples, figsize=(num_samples*4, 16))
def plot_sample(ax, idx, category):
    if len(X_test.shape) == 5:
        slice_idx = X_test.shape[1] // 2
        ax.imshow(X_test[idx, slice_idx, :, :, 0], cmap='gray')
    else:
        slice_idx = X_test.shape[1] // 2
        ax.imshow(X_test[idx, slice_idx], cmap='gray')
    prob = y_pred_prob[idx][0]
    ax.set_title(f'{category}\nProb: {prob:.2f}')
    ax.axis('off')
categories = ['True Positive', 'True Negative', 'False Positive', 'False Negative']
index_lists = [tp_indices, tn_indices, fp_indices, fn_indices]
for i, (category, indices) in enumerate(zip(categories, index_lists)):
    count = min(len(indices), num_samples)
    for j in range(count):
        plot_sample(axes[i, j], indices[j], category)
    for j in range(count, num_samples):
        axes[i, j].axis('off')
        axes[i, j].set_title(f'{category}\nN/A")
plt.tight_layout()
plt.savefig(os.path.join(save_dir, 'prediction_summary.png'))
plt.close()

```

```

def generate_feature_maps(model, sample, layer_names=None, save_dir=None):
    if not save_dir:
        save_dir = '3d_res/feature_maps'
    os.makedirs(save_dir, exist_ok=True)
    if not layer_names:
        layer_names = [layer.name for layer in model.layers if 'conv' in layer.name.lower()]
    feature_models = []

```

```

for layer_name in layer_names:

    feature_model = tf.keras.Model(inputs=model.input,
    outputs=model.get_layer(layer_name).output)

    feature_models.append((layer_name, feature_model))

if len(sample.shape) == 4:

    sample = np.expand_dims(sample, axis=0)

for layer_name, feature_model in feature_models:

    print(f"Generating feature maps for layer: {layer_name}")

    feature_maps = feature_model.predict(sample)

    if len(feature_maps.shape) == 5:

        n_features = feature_maps.shape[-1]

        middle_slice = feature_maps.shape[1] // 2

        grid_size = int(np.ceil(np.sqrt(min(n_features, 16))))

        fig, axes = plt.subplots(grid_size, grid_size, figsize=(12, 12))

        fig.suptitle(f"Feature Maps - Layer: {layer_name}", fontsize=16)

        feature_idx = 0

        for i in range(grid_size):

            for j in range(grid_size):

                if feature_idx < min(n_features, 16):

                    ax = axes[i, j] if grid_size > 1 else axes

                    im = ax.imshow(feature_maps[0, middle_slice, :, :, feature_idx], cmap='viridis')

                    ax.set_title(f"Feature {feature_idx+1}")

                    ax.axis('off')

                    feature_idx += 1

                else:

                    if grid_size > 1:

                        axes[i, j].axis('off')

                    else:

                        axes.axis('off')

        plt.tight_layout()

        plt.savefig(os.path.join(save_dir, f"feature_maps_{layer_name}.png"))

```

```

plt.close()

print(f'Feature maps saved to {save_dir}')

# Pipeline
def run_pipeline(base_dir, output_dir='3d_res'):
    print("="*80)
    print("LUNG CANCER DETECTION PIPELINE")
    print("="*80)

    # Create output directory
    os.makedirs(output_dir, exist_ok=True)

    # Step 1: Load preprocessed data
    data_dir = '.' # Try working directory first
    if not all(os.path.exists(os.path.join(data_dir, f)) for f in
               ['X_train.npy', 'X_val.npy', 'X_test.npy', 'y_train.npy', 'y_val.npy', 'y_test.npy']):
        data_dir = base_dir # Fall back to luna16_dataset
    print(f"Loading preprocessed data from {data_dir}...")
    try:
        X_train = np.load(os.path.join(data_dir, 'X_train.npy'))
        X_val = np.load(os.path.join(data_dir, 'X_val.npy'))
        X_test = np.load(os.path.join(data_dir, 'X_test.npy'))
        y_train = np.load(os.path.join(data_dir, 'y_train.npy'))
        y_val = np.load(os.path.join(data_dir, 'y_val.npy'))
        y_test = np.load(os.path.join(data_dir, 'y_test.npy'))
    except FileNotFoundError as e:
        raise FileNotFoundError(f"Preprocessed data not found in {data_dir}. Ensure X_train.npy, etc.,
are present: {str(e)}")

    # Add channel dimension if needed
    if len(X_train.shape) == 4:

```

```

X_train = np.expand_dims(X_train, axis=-1)
X_val = np.expand_dims(X_val, axis=-1)
X_test = np.expand_dims(X_test, axis=-1)

print(f"Data shapes:")
print(f" Training: {X_train.shape}, {y_train.shape}")
print(f" Validation: {X_val.shape}, {y_val.shape}")
print(f" Testing: {X_test.shape}, {y_test.shape}")

# Step 2: Load model
print("\nLoading model...")
try:
    model = tf.keras.models.load_model(
        os.path.join(data_dir, 'best_model.h5'),
        custom_objects={'focal_loss_fixed': focal_loss(gamma=2.0, alpha=0.25)}
    )
except FileNotFoundError:
    raise FileNotFoundError(f"Model file 'best_model.h5' not found in {data_dir}.")

# Step 3: Evaluate the model (skipped)
print("\nEvaluation already completed, proceeding to visualization...")

# Step 4: Analyze false negatives (skipped)
print("\nFalse negative analysis already completed, proceeding to visualization...")

# Step 5: Visualize Predictions
print("\nVisualizing model predictions...")
vis_dir = output_dir

# Get predictions
y_pred_prob = model.predict(X_test)

```

```
# Get indices for TP, TN, FP, FN
```

```
tp_indices = np.where((y_test == 1) & (y_pred_prob.flatten() > 0.5))[0]
```

```
tn_indices = np.where((y_test == 0) & (y_pred_prob.flatten() <= 0.5))[0]
```

```
fp_indices = np.where((y_test == 0) & (y_pred_prob.flatten() > 0.5))[0]
```

```
fn_indices = np.where((y_test == 1) & (y_pred_prob.flatten() <= 0.5))[0]
```

```
# 1. Visualize 3D volumes (animations for 2 positive and 2 negative samples)
```

```
print("Generating 3D volume animations...")
```

```
for idx in tp_indices[:2]:
```

```
    visualize_3d_volume(
```

```
        X_test[idx],
```

```
        title=f"Positive Sample (Nodule) - Index {idx}",
```

```
        save_path=os.path.join(vis_dir, 'true_positive', f"volume_positive_{idx}.gif")
```

```
    )
```

```
for idx in tn_indices[:2]:
```

```
    visualize_3d_volume(
```

```
        X_test[idx],
```

```
        title=f"Negative Sample (No Nodule) - Index {idx}",
```

```
        save_path=os.path.join(vis_dir, 'true_negative', f"volume_negative_{idx}.gif")
```

```
    )
```

```
# 2. Visualize prediction comparisons
```

```
print("Generating prediction comparison plots...")
```

```
categories = ['true_positive', 'true_negative', 'false_positive', 'false_negative']
```

```
index_lists = [tp_indices, tn_indices, fp_indices, fn_indices]
```

```
for category, indices in zip(categories, index_lists):
```

```
    count = min(len(indices), 3)
```

```
    for j in range(count):
```

```
        idx = indices[j]
```

```
        visualize_prediction_comparison(
```

```
    X_test[idx],
    y_test[idx],
    y_pred_prob[idx][0],
    save_dir=os.path.join(vis_dir, category)
)
```

3. Visualize model predictions (summary grid)

```
print("Generating prediction summary grid...")
```

```
visualize_model_predictions(
```

```
    model=model,
```

```
    X_test=X_test,
```

```
    y_test=y_test,
```

```
    num_samples=5,
```

```
    save_dir=vis_dir
```

```
)
```

4. Generate feature maps (for one positive and one negative sample)

```
print("Generating feature maps...")
```

```
if tp_indices.size > 0:
```

```
    generate_feature_maps(
```

```
        model,
```

```
        X_test[tp_indices[0]],
```

```
        save_dir=os.path.join(vis_dir, 'feature_maps_positive')
```

```
    )
```

```
if tn_indices.size > 0:
```

```
    generate_feature_maps(
```

```
        model,
```

```
        X_test[tn_indices[0]],
```

```
        save_dir=os.path.join(vis_dir, 'feature_maps_negative')
```

```
    )
```

```
print(f"Visualizations saved in {vis_dir}")
```

```
print("\nPipeline completed successfully!")
```

```
print(f"Results saved to {output_dir}")
```

```
return model, {}
```

```
if __name__ == "__main__":
```

```
    base_dir = '/workspace/deep_learning_luna16/luna16_dataset'
```

```
    model, metrics = run_pipeline(base_dir)
```

OUTPUT OF BLOCK 3

```
=====
LUNG CANCER DETECTION PIPELINE
=====
Loading preprocessed data from ....
Data shapes:
  Training: (2527, 64, 64, 64, 1), (2527,)
  Validation: (542, 64, 64, 64, 1), (542,)
  Testing: (542, 64, 64, 64, 1), (542,)

Loading model...

Evaluation already completed, proceeding to visualization...

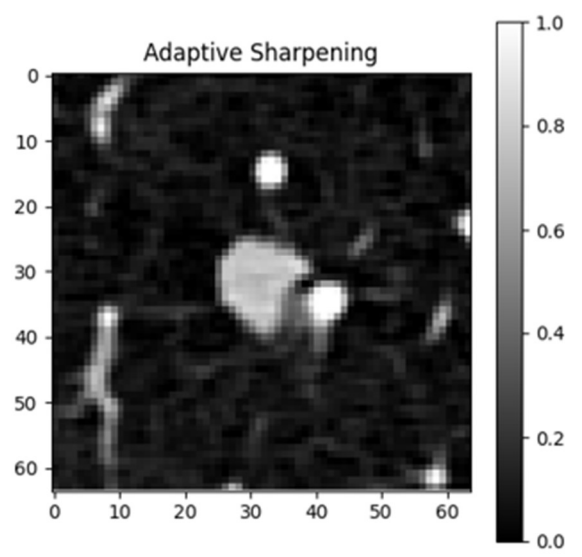
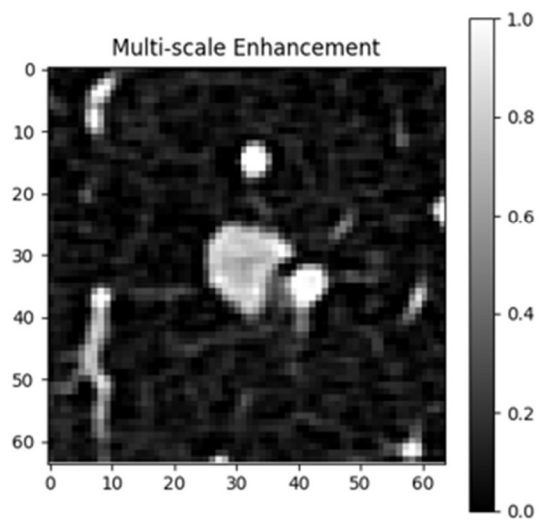

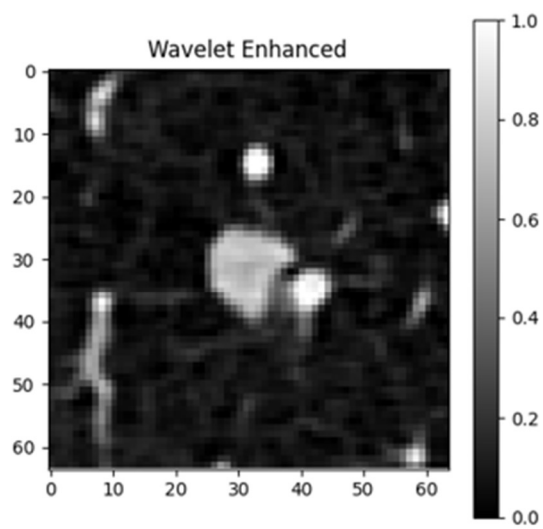
False negative analysis already completed, proceeding to visualization...

Visualizing model predictions...
17/17 [=====] - 2s 92ms/step
Generating 3D volume animations...
Generating prediction comparison plots...
Generating prediction summary grid...
17/17 [=====] - 1s 75ms/step
Generating feature maps...
Generating feature maps for layer: conv3d_7
1/1 [=====] - 0s 159ms/step
Generating feature maps for layer: conv3d_8
...
Visualizations saved in 3d_res


Pipeline completed successfully!
Results saved to 3d_res
```

OUTPUTS(RESULTS)


- Enhancement:




- Prediction Summary




- True positive sample:



- True Negative sample:



- False Positive



- False Negative:

